

# Concurrency and Parallel Programming

## Test results on DAS-4

Maarten de Jonge

November 7, 2012

### Experimental setup

The wave equation has been straightforwardly implemented. When  $n_p$  points on the wave are being simulated on  $n_t$  threads, each thread handles a slice of roughly  $\frac{n_p}{n_t}$  points, with some variation due to the division not being guaranteed to be exact.

The simulation has been run on the DAS-4 computer with various values for each parameter; The number of simulated points is tested at 1.000, 10.000, 1.000.000, 10.000.000 (with dots added to increase readability), and each of the values has been tested with 1, 2, 4, 8 and 16 threads. In each case, a suitable value for the number of time steps has been selected such that the total time taken with 1 thread is somewhere around 50 seconds, to allow for decently accurate benchmarking without getting bored. These values can be seen in table 1.

number of points	number of iterations
$10^3$	1e6
$10^4$	5e5
$10^6$	5e3
$10^7$	5e2

Table 1: The number of points with their chosen number of timesteps

### Results

Figure 1 shows the results. When simulating 1000 or 10000 points, there is an initial speedup when increasing the amount of threads, but after a certain point any additional threads will dramatically lower performance. The higher amounts of points scale very well with the amount of threads, up to 8 threads (which is the amount of cores available on system). Scaling it up to 16 threads does not noticeably improve performance, although it doesn't degrade it either.

### Conclusion

There appears to be some overhead to the use of threads; they can offer a significant speedup for the tested scenario, but only when each thread has a




Figure 1: The test results. "i\_max" refers to the amount of points simulated on the wave.

i_max	t_max	num_threads	time	normalized
1000	1000000	1	42.4294	$4.24294e-08$
1000	1000000	2	31.5108	$3.15108e-08$
1000	1000000	4	80.9137	$8.09137e-08$
1000	1000000	8	167.284	$1.67284e-07$
1000	1000000	16	350.395	$3.50395e-07$
10000	500000	1	69.3228	$9.83256e-08$
10000	500000	2	44.066	$6.2502e-08$
10000	500000	4	41.1647	$5.83869e-08$
10000	500000	8	85.3845	$1.21107e-07$
10000	500000	16	183.268	$2.59943e-07$
1000000	5000	1	52.108	$7.39086e-08$
1000000	5000	2	25.6232	$3.63433e-08$
1000000	5000	4	13.1806	$1.86951e-08$
1000000	5000	8	7.08639	$1.00511e-08$
1000000	5000	16	7.06076	$1.00148e-08$
10000000	500	1	51.4277	$7.29436e-08$
10000000	500	2	25.8168	$3.66179e-08$
10000000	500	4	13.0537	$1.85151e-08$
10000000	500	8	7.21982	$1.02404e-08$
10000000	500	16	7.16489	$1.01625e-08$

Table 2: The raw test data, where “i\_max” is the number of simulated points on the wave and “t\_max” is the amount of iterations.

high enough workload. It seems that when there’s too little data to simulate per thread, the work done by the thread will not be worth the incurred overhead, leading to a (quite severe) degradation of performance.